

ASSIGNMENT OF MASTER'S THESIS

Title: Improvements of the RIR bytecode toolchain
Student: Bc. Jan Je men
Supervisor: Ing. Petr Máj
Study Programme: Informatics
Study Branch: System Programming
Department: Department of Theoretical Computer Science
Validity: Until the end of winter semester 2018/19

Instructions

Familiarize yourself with the R language, its bytecode compiler, and interpreter architecture. Familiarize yourself with RIR, an alternative bytecode format, compiler, and interpreter for the language. The R bytecode compiler assumes certain invariants (such as built-in meaning of control flow statements and certain operators) about the code to make the compiled code faster. Analyze similar assumptions that are used by RIR and extend RIR to use assumptions made by GNU-R as well. Identify and implement improvements to the RIR (compiler, bytecode format, and interpreter). Discuss your results.

References

Will be provided by the supervisor.

doc. Ing. Jan Janoušek, Ph.D.
Head of Department

prof. Ing. Pavel Tvrdík, CSc.
Dean

Prague March 1, 2017

CZECH TECHNICAL UNIVERSITY IN PRAGUE

FACULTY OF INFORMATION TECHNOLOGY

DEPARTMENT OF THEORETICAL COMPUTER
SCIENCE



Master's thesis

Improvements of the RIR bytecode toolchain

Bc. Jan Ječmen

Supervisor: Ing. Petr Máj

May 4, 2017

Acknowledgements

[[acknowledgements]] Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Declaration

I hereby declare that the presented thesis is my own work and that I have cited all sources of information in accordance with the Guideline for adhering to ethical principles when elaborating an academic final thesis.

I acknowledge that my thesis is subject to the rights and obligations stipulated by the Act No. 121/2000 Coll., the Copyright Act, as amended, in particular that the Czech Technical University in Prague has the right to conclude a license agreement on the utilization of this thesis as a school work under the provisions of Article 60(1) of the Act.

In Prague on May 4, 2017

.....

Czech Technical University in Prague

Faculty of Information Technology

© 2017 Jan Ječmen. All rights reserved.

This thesis is a school work as defined by Copyright Act of the Czech Republic. It has been submitted at Czech Technical University in Prague, Faculty of Information Technology. The thesis is protected by the Copyright Act and its usage without author's permission is prohibited (with exceptions defined by the Copyright Act).

Citation of this thesis

JEČMEN, Jan. *Improvements of the RIR bytecode toolchain*. Master's thesis. Czech Technical University in Prague, Faculty of Information Technology, 2017. Available also from WWW: [⟨https://github.com/JanJecmen/dip⟩](https://github.com/JanJecmen/dip).

Abstract

[[abstract english]] This is the second paragraph. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Keywords R language, RIR, bytecode compiler & interpreter, optimizations

Abstrakt

[[abstract czech]] And after the second paragraph follows the third paragraph. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Klíčová slova jazyk R, RIR, kompilátor & interpret bajtkódu, optimalizace

Contents

Introduction	1
1 About GNU R	3
1.1 Language features of R	4
1.2 AST interpreter	11
1.3 BC interpreter and compiler	15
1.4 Why is R hard to optimize	21
2 About RIR	23
2.1 Why is RIR slow	24
3 Improvements	25
3.1 Instruction set extensions	25
3.2 Compiler modifications	30
3.3 Interpreter refactoring	31
4 Evaluation	33
Conclusion	37
Bibliography	39
A Acronyms	41
B Contents of the enclosed CD	43

List of Figures

4.1	History of running times	34
4.2	Overview of the speedup vs. GNU R	35
4.3	History of average speedup vs. GNU R	36

List of Listings

1.1	Anonymous function	4
1.2	Arithmetic operators are function calls in R	5
1.3	Promise lazy evaluation	6
1.4	Superassignment	6
1.5	Coercion to the most flexible type	7
1.6	Recycling shorter vector	8
1.7	Object attributes	8
1.8	Argument handling	9
1.9	S3 object system	10
1.10	Basic subsetting / subassignment	11
1.11	Other subsetting operators	12
1.12	AST of a simple expression	14
1.13	Disassembling a BC object	17
1.14	Breaking the compiler	20
1.15	Reason for the erroneous result	21
3.1	Complex subset assignment	29
3.2	Safe break	30
3.3	Context for break required	30
3.4	31

List of Tables

1.1	Some common types of internal R objects	13
1.2	Description of some GNU R bytecodes	17

Introduction

[[write intro; cite:
<https://www.tiobe.com/tiobe-index/>
<http://pypl.github.io/PYPL.html>]]

This thesis is structured in the following way:

The chapter *About GNU R* gives a short introduction to GNU R and discusses its features. It also goes under the hood and describes the inner workings of its interpreter and bytecode compiler.

The chapter *About RIR* introduces an alternative bytecode compiler for the R language, talks about the motivation behind it, its architecture and design choices, the differences to the original and its shortcomings.

The chapter *Improvements* describes in depth the changes that were done to RIR in this thesis. [[?]]

The chapter *Evaluation* discusses how the performance of RIR changed after the changes done in this thesis. It describes how the measurements were done and how the performance compares to GNU R. [[?]]

The results of this thesis are discussed in *Conclusion*, as well as the direction of future efforts regarding RIR.

About GNU R

[[cite]] GNU R is a programming language used mainly for statistical computations.¹ It is an open-source dialect of S, an older statistical language created in 1976 by John Chambers at Bell Laboratories. R has been around from 1993 and was designed by Ross Ihaka and Robert Gentleman, both recognised statisticians. It is a part of the GNU software family and is still actively developed by the R Core Team today. It is a popular alternative to the other major implementation of the S language, S-PLUS, which is a commercial version shipped by TIBCO Software Inc.

R comes with a software environment built around it, which allows for easily manipulating data, carrying out computations and producing quality graphical outputs such as plots and figures. Although at heart R is used via a command line interface, there are also more user-friendly graphical IDEs available. One of the most widely used is RStudio that provides, for example, syntax highlighting and quick access to documentation through a web-like browser.² This, together with R's readable syntax and a vast collection of extension packages available through The Comprehensive R Archive Network (CRAN) makes it possible for new users to step in and start working quickly.

¹Homepage: <https://www.r-project.org/>

²Homepage: <https://www.rstudio.com/>

1.1 Language features of R

[[cite: <http://r.cs.purdue.edu/pub/ecoop12.pdf>, <http://adv-r.had.co.nz/> ...]] R is, as far as programming languages go, very interesting and has some quite unusual semantic features. It is an interpreted language, and is dynamically typed and garbage collected. It supports multiple programming paradigms: users can use procedural imperative style, but at the same time R provides an object system (more than one, in fact!) for object oriented programming, and is heavily influenced by functional programming languages, notably Scheme (a dialect of Lisp).

Functions are, in accordance with functional languages, first-class values, so they can be passed around as call arguments, returned as results of function calls and created dynamically at runtime. R uses lexical scoping (which it adopted from Scheme) and R functions are closures that capture their enclosing environment at creation time. Arguments are passed by value (although a variant of reference counting is implemented, so that deep copies are only created as needed, e.g., when an object is modified).

Functions are by design anonymous, i.e., they are not named when created. This is unlike many languages (e.g., C or Python) and follows the approach of lambda calculus. Instead of creating named functions, R programmers create anonymous ones and, if they choose so, then use regular assignment to bind them to names. Example can be seen in listing 1.1.

```
> (function(x) x + 1)(2)
[1] 3
> f <- function(x) x + 1
> f(2)
[1] 3
```

Listing 1.1: Anonymous function

Interestingly, everything that happens in R is in fact a function call. This goes as far as arithmetic operators being just syntactic sugar for function calls, as can be seen in listing 1.2.³ In this spirit, even assigning into a variable, evaluating

³As a side note, backticks are used in R to denote symbols (i.e. names). Because some symbols have special syntactic meaning (like the ``+`` being an infix binary operator), they can cause a syntax error if they appear unquoted in a place where the parser does not expect them.

a block of code inside curly brackets or grouping expressions with parentheses translate to calling the respective functions.

```
> typeof('+')
[1] "builtin"
> '+'
function (e1, e2) .Primitive("+")
> '+'(1, 2)
[1] 3
> 1 + 2
[1] 3
```

Listing 1.2: Arithmetic operators are function calls in R

All actual arguments to a function are lazy evaluated by default. When applying a closure, parameters are wrapped in promises. A promise is simply an object that contains the unevaluated expression and an environment in which the expression should be evaluated. Promises are only evaluated when the value is actually needed (which is called forcing the promise). Also, once the promise is forced, it remembers the result, so that subsequent uses of the value do not evaluate the expression again.

Delayed evaluation is demonstrated in listing 1.3, where the function in question only evaluates its first argument and does not touch the second. For the first call, the block of code in the curly brackets is never executed and so the side-effect of printing a greeting does not happen.⁴ In the second call, the arguments are swapped, and the side-effect can be observed in the output. It is possible to pass a block of code as a function argument, since the `{ }` is bound to a function that sequentially evaluates all its arguments and returns as its result the value of the last expression (or `NULL` if no expressions are passed in).

These features highlight the functional approach of R by minimizing side-effects. However, R supports assignment which enables programmers to change a function's local state by modifying its bindings and thus the imperative programming style.⁵ Also, the superassignment operator `<<-` makes it possible to

⁴`cat` is short for *concatenate and print*.

⁵An interesting quirk is R's left to right assignment: the code `1 -> x` assigns to `x` and is equivalent to `x <- 1`. It has nothing to do with C++ structure dereference, since R does not have references.

```
> f <- function(a, b) a
> f(1, {cat("Hello\n"); 2})
[1] 1
> f({cat("Hello\n"); 2}, 1)
Hello
[1] 2
```

Listing 1.3: Promise lazy evaluation

change non-local bindings (as shown in listing 1.4) and thus brings the side-effects back into play.

```
> x <- 1
> f <- function() {
+   x <- 2 # local
+   x <- 3 # lookup in the enclosing environment
+   x # local
+ }
> x
[1] 1
> f()
[1] 2
> x
[1] 3
```

Listing 1.4: Superassignment

The basic data type in R is an atomic vector. Vectors are collections of homogeneous values (a given vector can only hold objects of one particular type), that preserve the order of their elements. R also provides a list type which is heterogeneous. Higher-dimensional types such as matrices and data frames, as well as objects, are built from vectors and lists under the hood. Vectors can be created in R by calling the function `c`.⁶

Atomic vectors can have one of these six types: logical, integer, double, character, complex and raw. Since R targets data analysis, a special “not available” value `NA` is provided for these. Because all values in a vector must be of the same type, R performs coercion when an attempt is made to combine vectors of different types. In listing 1.5, combining a numeric vector with a character

⁶The name stands for *combine*.

vector results in a character vector, and summing a logical vector coerces to integers (**TRUE** becoming 1 and **FALSE** becoming 0).

In R there are no scalar types, as scalar values, such as individual numbers and strings, are considered to be vectors of length one. This holds for character vectors, too, which can cause confusion if one expects C-like behavior of strings and tries subscripting, because in R the subscript belongs to the vector that holds the string and not the string itself.⁷

```
> x <- c(1, 2, 3)
> y <- c("a", "b", "c")
> typeof(x)
[1] "double"
> typeof(y)
[1] "character"
> c(x, y)
[1] "1" "2" "3" "a" "b" "c"
> typeof(c(x, y))
[1] "character"
>
> sum(c(TRUE, TRUE, FALSE, TRUE))
[1] 3
```

Listing 1.5: Coercion to the most flexible type

Despite its inspiration in functional world, R does not optimize tail recursion, which is the standard approach in functional languages since they typically use recursion in the place of iterative loops. Instead, R encourages vectorized operations. Hence, most of R builtin functionality works element-wise with vectors, while recycling the elements as needed (e.g., when adding vectors of different lengths). This is demonstrated in listing 1.6, where R even issues a warning about recycling.

In R, every object can have arbitrary attributes associated with its data. Attributes are basically a hidden map that assigns names to values. Some of the most important attributes are names (a character vector that assigns names to elements), dimensions (a vector specifying the dimensions and thus effectively distinguishing vectors from matrices and arrays) and class (for implementing one of R's object systems). Attributes are used a lot in R for many purposes and

⁷In such a case, one needs to use the **substr** function.

```
> numeric(10)
[1] 0 0 0 0 0 0 0 0 0 0
> 1:3
[1] 1 2 3
> numeric(10) + 1:3
[1] 1 2 3 1 2 3 1 2 3 1
Warning message:
In numeric(10) + 1:3 :
  longer object length is not a multiple of shorter object length
```

Listing 1.6: Recycling shorter vector

extensions as they provide a way of encoding arbitrary additional metadata for objects. As an example, in listing 1.7 a vector of 4 elements is created, then changed into a two by two matrix and finally changed back to vector (that also has its elements named).⁸

```
> x <- 1:4
> x
[1] 1 2 3 4
> attr(x, "dim") <- c(2, 2)
> x
      [,1] [,2]
[1,]    1    3
[2,]    2    4
> attr(x, "dim") <- c(4)
> attr(x, "names") <- c("a", "b", "c", "d")
> x
a b c d
1 2 3 4
```

Listing 1.7: Object attributes

True to its dynamic nature, R is very liberal in handling arguments in function calls. The language supports both positional and named matching of arguments, as well as default argument values. R even understands when the argument names are abbreviated, as long as the arguments can still be uniquely matched.

Moreover, R lets users call functions and not provide the specified arguments. It is only while executing the function body that the missing arguments may

⁸Here it can be seen here that R uses column-major order for storing the elements.

or may not cause an error. Variable number of arguments is supported as well by using the ellipsis ``...``. The ellipsis in an argument list matches any number of arguments, and later in the function body refers to the list of matched arguments. Special symbols can be used to access the ellipsis arguments, such as ``..1``, ``..2``, etc. Argument handling is demonstrated in listing 1.8.

```
> f <- function(a, b, a.very.long.argument.name)
+   a.very.long.argument.name
> f(1, 2, 3)
[1] 3
> f(a = 3)
Error in f(a = 3) :
  argument "a.very.long.argument.name" is missing, with no default
> f(a. = 3)
[1] 3
>
> f <- function(a, b) a
> # b is not required
> f(1)
[1] 1
> # a is required
> f()
Error in f() : argument "a" is missing, with no default
> f(b = 2)
Error in f(b = 2) : argument "a" is missing, with no default
>
> f <- function(...) ..2
> f()
Error in f() : the ... list does not contain 2 elements
> f(1, 2, 3, 4, 5)
[1] 2
```

Listing 1.8: Argument handling

As was already mentioned, R has not one but three different object systems. The simplest is called S3 and it uses a class attribute to implement ad hoc polymorphism (also known as function or operator overloading). It does not have formal classes, but instead uses a special function called a generic function that decides what to call based on the value of the class attribute. A typical example is printing, where the `print` function is generic, its body consists of a call to a dispatcher `UseMethod("print")`. Specialized versions for different types can be defined, such as `print.data.frame` for data frames, or, given an object with class set to `"foo"`, `print.foo` (as is shown in listing 1.9).

```
> x <- list()
> class(x) <- "foo"
> print(x)
list()
attr(,"class")
[1] "foo"
> print.foo <- function(...) cat("Printing foo!\n")
> print(x)
Printing foo!
```

Listing 1.9: S3 object system

S4 is a more formal system than S3, and it allows for true class definitions, describing class representation and inheritance. It has multiple dispatch, which means that dispatchers can pick which method to call based on multiple arguments. R also has reference classes that implement message passing style. Their objects are modified in place (as opposed to the standard pass by value semantics).

An important and powerful feature of R is its subsetting and subset assignment mechanics. Parts of objects can be retrieved and even changed by the operators ``[``, ``[[`` and ``$``. These behave differently for different types of objects.

The first version is a general subset function that supports all kinds of indexing.⁹ For example, integer vector specifies which elements to get and in what order, even allowing duplication. If negative indices are used, these elements are omitted from the result. Logical vectors can be used to select only elements at positions where `TRUE` occurs. If the object is named, character vectors can be used to select by name. In combination with assignment (and superassignment), objects can be modified. Subsetting works also for higher-dimensional structures by simply providing indices for each dimension, separated by commas. Some examples are in listing 1.10.

The second version, ``[[``, returns only a single element of an object, and is used to get elements out of a list. The ``$`` is then just a shorthand for ``[[`` when used with character subsetting (that is, the dollar version expects a name, and it need not be quoted). Also, the dollar operator does partial matching, similar to

⁹As opposed to many languages (e.g., C and Python), R starts indexing elements from 1 instead of 0.

```

> x <- 101:110
> x
[1] 101 102 103 104 105 106 107 108 109 110
> x[c(3, 5, 1, 1)]
[1] 103 105 101 101
> x[-c(2, 3, 8)]
[1] 101 104 105 106 107 109 110
> x > 105
[1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
> x[x > 105]
[1] 106 107 108 109 110
> x[x %% 2 == 0] <- NA
> x
[1] 101 NA 103 NA 105 NA 107 NA 109 NA
> m <- matrix(1:9, ncol = 3)
> m
      [,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
> m[-2, 2:3] # omit row 2 and get columns 2 to 3
      [,1] [,2]
[1,]    4    7
[2,]    6    9

```

Listing 1.10: Basic subsetting / subassignment

how function argument names are handled. These operators are demonstrated in listing 1.11.

1.2 AST interpreter

In its core R uses a classic architecture for an interpreted language. After initialization, the user enters R's read-eval-print loop (REPL), that lets them type in expressions and have R evaluate them. First, a reader, or parser, waits for the user input and reads it line by line. If, at the end of line, it has read a syntactically complete expression, it passes it to an evaluator (otherwise it waits for more input). After the evaluator returns the evaluated expression, a printer is invoked that displays the result (with some exceptions, such as assignment that sets its result to be invisible). Then the reader is invoked again and the process repeats.

```
> l <- list(sq = 1:3, str = "a", bool = FALSE, na = NA)
> l
$sq
[1] 1 2 3

$str
[1] "a"

$bool
[1] FALSE

$na
[1] NA

> l[1]
$sq
[1] 1 2 3

> typeof(l[1])
[1] "list"
> l[[1]]
[1] 1 2 3
> typeof(l[[1]])
[1] "integer"
> l[["bool"]]
[1] FALSE
> l$bool
[1] FALSE
```

Listing 1.11: Other subsetting operators

Every object in R is internally represented by a C structure called SEXPREC (actually, R passes the objects around as pointers to this structure, which are called SEXP).¹⁰ This structure contains a header with metadata about the object, such as its type, reference counter or information for garbage collector, and then a union of other structures that represent different types of R internal objects. Some of these types are listed in table 1.1.

¹⁰The name refers to S-expressions, or symbolic expressions, as known from Lisp, although the classical linked lists built from dotted pairs are mostly used internally, and vectors are implemented as C arrays for efficiency reasons.

Table 1.1: Some common types of internal R objects

Type	Usage
NILSXP	the singleton NULL object
SYMSXP	symbols (or names)
LISTSXP	lists of dotted pairs
CLOSXP	closures
ENVSXP	environments
PROMSXP	promises
LANGSXP	language constructs (typically closure application)
SPECIALSXP	special forms (typically control flow)
BUILTINSXP	builtin non-special forms (e.g., arithmetic operators)
INTSXP	integer vectors
REALSXP	real vectors
STRSXP	string vectors
BCODESXP	object compiled to bytecode

The parser, when it scans the stream of input characters, checks that it is syntactically correct and at the same time builds a tree structure that represents the parsed expression. This tree is called the abstract syntax tree (AST) and its nodes are all SEXPs. An example AST is shown in the listing 1.12.¹¹ In the listing, parentheses denote function calls (i.e. LANGSXP node), the first child being the callee (typically a SYMSXP, i.e., a name that is bound to a function) and the rest its arguments.

The evaluator is a recursive function that gets as its input two SEXP objects, one representing the AST of the expression that is to be evaluated, and the other the environment in which to evaluate the expression. The evaluator walks the given AST and based on the type of nodes it encounters, performs some action. The result is then returned to be processed by the caller of eval.

¹¹pryr (homepage: <https://github.com/hadley/pryr>) is a package created by Hadley Wickham that allows to “pry back the surface of R and dig into the details.”

```
> pryr::ast(x <- (y + 3) * f(z))
\ - (
  \ - `<-
    \ - `x
      \ - (
        \ - `*
          \ - (
            \ - `(
              \ - (
                \ - `+
                  \ - `y
                    \ - 3
              \ - (
                \ - `f
                  \ - `z
```

Listing 1.12: AST of a simple expression

Some nodes are self-evaluating, meaning that no action needs to be performed and the node itself is the result. These are for example the **NULL** object, atomic vectors or environments.

If the eval function sees a symbol node, a lookup for its binding is performed in the provided environment. If it is not found there, because of lexical scoping, the parent environment is searched, and so on, until either the binding is found or an empty environment is reached (the empty environment serves as a sentinel parent of all environment chains and does not have a parent itself).

One other prominent type of nodes is LANGSXP. R has internally three types of functions, called *special*, *builtin* and *closures* (or user-defined functions). These have different behavior when they are applied, and the eval function handles that.

Special functions are the core language constructs, such as control flow (conditionals and loops). They take their arguments unevaluated in a list and evaluate them as needed while running. This is necessary for example for the **if** statement, because, since R has side-effects, only one of the conditional branches must be evaluated to preserve the correct semantics.

Builtins, on the other hand, are known to evaluate all their arguments, so it is not necessary to create promises from their arguments. Instead, a list of

evaluated arguments is created and passed to the builtin function. Examples of builtin functions are arithmetic operators or the colon operator for generating integer sequences.

The last group are closures. Closures are user-defined functions written in R, and they adhere to the lazy evaluation semantics. All arguments to a closure are therefore allocated as promises: the expressions are bundled together with the enclosing environment.

As opposed to specials and builtins, which, being C routines, are called directly after preparing their arguments, closures need the interpreter to do some additional work. First, the actual unevaluated arguments have to be matched to the formal arguments of the closure. Then, a new environment has to be created and filled with the matched pairs of arguments. Only after this can the body of the closure be evaluated in the new environment. Also, a longjump target is set here to catch any explicit return calls from within the body.

Finally, the dispatch to the bytecode interpreter for objects compiled to bytecode is also found in the eval function.

1.3 BC interpreter and compiler

In an attempt to make R faster, a special internal representation for R code was developed, and a compiler from R to this bytecode was added in a package called `compiler`. This also required some minor changes to the original AST interpreter, namely adding a new `SEXP` type for the compiled objects, called `BCODESEXP`, and handling of bytecode objects in the evaluator. A new evaluator was also added for interpreting the bytecode which is invoked by the AST version when it needs to evaluate a compiled object.

The compiler was written by Luke Tierney, and added as a standard package to R in 2011 in version 2.13.0. However, it was not used by default until version 3.4.0, released in late April 2017 ([2] and [3]).¹²

[[cite compiler pdf]] The BC compiler itself is implemented in R, and walks the abstract syntax tree of an expression being compiled in a similar manner

¹²The default packages were compiled, but for additional packages and user code JIT was disabled and had to be explicitly enabled.

that the AST interpreter does (but, of course, it does so at the R level by using introspection). However, instead of evaluating the code as it traverses the tree, it produces a code object. The code is then later executed by the BC interpreter, a separate virtual machine runtime system from the AST version. The compiler uses just a single pass, meaning that it only looks at the compiled expression once, and while doing so, produces a stream of instruction. A multi-pass version that would add optimization passes for the internal representation is planned to be explored in the future.

The bytecode objects produced by the compiler consist of two components. The first is an integer vector that encodes the code itself in the form of instruction opcodes interleaved with the instruction operands. The second is a general list that represents a constant pool. In the constant pool, important objects are stored, such as the source for the compiled expression, small constant objects, or promises. The compiler is designed such that each bytecode object has its own constant pool.

The compiler can be used explicitly to compile an expression or a closure. However, a more convenient way is to enable just-in-time compilation (JIT). Doing so causes the AST interpreter to invoke the compiler automatically when calling a closure that is not yet compiled.

The compiler comes with a disassembler that makes it possible to inspect the bytecode of an object, as is shown in listing 1.13.¹³ The object is printed as a list that starts with the `.Code` symbol, then follows the code vector (with the opcodes decoded), and last comes the constant pool.

The integers in the code that are not instructions represent immediate arguments to the instructions (the first element being an exception, as it encodes the version of the BC stored in the given object). In this particular instance, all immediates represent constant pool indices.

The virtual machine that executes R bytecode uses a stack oriented architecture. This means that a stack is used by the instructions at runtime to get their arguments and store their results. For example, the instruction that performs addition, expects its two operands at the top of the stack. When it is executed,

¹³The output was reformatted for the sake of readability.

```

> f <- compiler::cmpfun(function(n) n + 1)
> f
function(n) n + 1
<bytecode: 0x367ee40>
> compiler::disassemble(f)
list(.Code,
     list(8L,
          GETVAR.OP, 1L,
          LDCONST.OP, 2L,
          ADD.OP, 0L,
          RETURN.OP),
     list(n + 1,
          n,
          1)
)

```

Listing 1.13: Disassembling a BC object

it removes these two objects from the stack, adds them together, and puts the resulting object back on the top of the stack.

The VM is implemented as a C routine that gets a bytecode object and an environment as arguments (similar to the AST interpreter). It verifies the BC version and then enters a loop that looks at the instruction stream in the BC object and dispatches to code that implements the given instruction. The loop is very carefully optimized by various techniques, such as using C preprocessor macros and threaded code. This will be discussed later in chapter 3.

The instruction set of the internal representation is designed to allow big parts of the AST interpreter internals to be reused. There are currently 123 instructions, some of which are described in table 1.2.

Table 1.2: Description of some GNU R bytecodes

Instruction	Description
RETURN.OP	Take the top of stack and return it as a result
GOTO.OP	Unconditionally jump to a label
BRIFNOT.OP	Conditionally jump to a label
POP.OP	Remove the top of stack value
LDCONST.OP	Push a constant from the constant pool
GETVAR.OP	Look up the symbol binding and push it

1. ABOUT GNU R

Instruction	Description
SETVAR.OP	Update the symbol binding
MAKEPROM.OP	Create promise from a call argument
CALL.OP	Do function call
CALLBUILTIN.OP	Call builtin function
CALLSPECIAL.OP	Call special function
MAKECLOSURE.OP	Create closure (with environment)
ADD.OP	Arithmetic binary plus
LT.OP	Relational less than
STARTASSIGN.OP	Prepare for subassignment
ENDASSIGN.OP	Clean up after subassignment
ISNULL.OP	Test if top of stack is NULL
COLON.OP	Create integer sequence

The compiler itself has in its heart the recursive function `cmp` that visits the AST of a given expression. It passes along a code buffer object (that contains the instruction stream and the constant pool) and a context object. When generating code, it writes into the code buffer, and uses the context to guide the compilation (it carries along information such as whether the expression should be followed by a return, if the result is ignored or not or if the expression is in a loop).

The expressions that are not self-evaluating are function calls, variable references, bytecode objects and promises. The rest is treated as being a constant. Bytecode objects and promises should not appear as literals in code, so they cause a compilation error if encountered.

Constant expressions are compiled by inserting the object into constant pool and generating a load instruction such as `LDCONST.OP` (that takes as an argument the index into the constant pool of the object).

For variable references, the symbol is inserted into the constant pool and then a `GETVAR.OP` instruction is emitted (although there is a special instruction for the “dot-dot-names” such as `..1`).

Everything else is a function call. When compiling a function call, multiple steps are required. First, the function to call has to be compiled. Usually this involves emitting an instruction that looks up the function by its name, but sometimes also compiling an expression that evaluates to a function. Then the arguments are compiled and code that prepares them on the stack is emitted. Finally, the call instruction is generated.

Since the compiler uses only a single pass, it has limited options of optimizing the generated code. The only optimization it performs, apart from those described in the next section, is constant folding. This is a very useful transformation that attempts to replace subtrees of an expression's AST that are constant (not only self-evaluating, but rather always evaluating to the same result) with the constant result.

Currently, constant folding is performed (depending on compiler options) on “small” expressions that consist of self-evaluating expressions, select base variables (like `pi`) and calls to some select base math functions (like `sqrt` or `sin`). In the current version, no deoptimization is possible.

1.3.1 BC compiler assumptions

When dealing with a language as dynamic as R, one needs to be very careful to distinguish compile time from execution time, and in particular to realize that between the two there is a window during which the state of the program can change.

It was explained earlier (see listing 1.2) that everything that happens in R is a function call, even constructs that have a special syntactic meaning, like `if`, `while`, `break` etc.

Furthermore, R uses lexical scoping. When the user starts an interactive session, they work in the global environment, which is the lowest environment in the hierarchy of environments known as the search path. The base environment (where all the default R bindings are defined) as well as any loaded packages are its ancestors in the search path. It is therefore possible to shadow the default bindings by creating new bindings that are found earlier on the search path.

Taking all this into consideration, to be conservative the compiler would have to compile everything as a function call, and then, at runtime, look up the current value of the callee binding and perform the call. Unfortunately, this almost completely negates the benefits of having a compiler at all, because the call is exactly what the AST interpreter does.

For this reason, the compiler assumes about certain functions that their meaning does not change in the meantime between compilation and execution. This is quite a reasonable assumption, since bindings for these special and builtin functions are rarely redefined, and it allows them to be properly compiled.

As a downside, the compiler produces code that is incorrect if the assumptions do not hold. Moreover, the bytecode interpreter has no way of telling that the code became incorrect, so it just executes it without any warning or error. To demonstrate, consider the code in listing 1.14.¹⁴

```
> f <- function(n) n + 1
> fc <- compiler::cmpfun(f)
> c(f(1), fc(1))
[1] 2 2
> `+` <- `-`
> c(f(1), fc(1))
[1] 0 2
```

Listing 1.14: Breaking the compiler

The reason for this behavior can be seen in listing 1.15 which shows the relevant part of the disassembled `fc` function. The ``+`` was inlined at compile time (because at the time, it really was an addition). When executing the function for the second time, however, the meaning of ``+`` was actually subtraction, and this was not reflected in the bytecode.

To reduce the chance of such errors, the compiler keeps track of local and global bindings visible at compile time. If a shadowing is detected the function in question is not inlined.

¹⁴JIT needs to be turned off for this.


```
list(8L,  
     GETVAR.OP, 1L,  
     LDCONST.OP, 2L,  
     ADD.OP, 0L,  
     RETURN.OP)
```

Listing 1.15: Reason for the erroneous result

1.4 Why is R hard to optimize

[[TODO]][1] [4] [7] [[cite: <https://cran.r-project.org/doc/manuals/R-lang.html> section 6 and 2, hadley]]

Firstly, R is not the fastest language out there. Of course, being an interpreted language, one cannot expect the performance of languages like C that are compiled to native machine code. This is because during runtime, there is the inherent overhead of managing the virtual machine that executes a given program. For the AST interpreter, it entails walking the tree again with every evaluation of an expression. For the BC compiler and interpreter, there is first the compilation itself (which only happens once), but then dispatching of the instructions needs to be done in the interpreter loop.

Another matter is that R is inherently single threaded, it is very memory hungry (as it needs a lot of metadata about its internal structures) and all memory is allocated on the heap and managed and garbage collected by the runtime system of R.

Just to give an example, incrementing a variable in C is done in one machine instruction (and possibly one memory store if the variable is not allocated in a register). An AST interpreter has to navigate a tree data structure in memory which, for this example, would probably consist of at least one assign node, two variable lookup nodes, one constant node and one arithmetic operation node. For the bytecode interpreter, the amount of work is considerably reduced, but still there is the large gap between machine instructions and the same bytecode instructions executed by a virtual machine (that must, for example, perform all type checking and possibly type promotions dynamically at runtime).

On top of that, R is a very dynamic language that gives a very high degree of freedom to a programmer. The design decisions behind R have an impact on performance. At runtime, users have full access to all of the program data and representation. This means, for example, that not only the values of a function's arguments can be accessed, but also the code that is used to compute them.

Even though R did not go as far as Lisp which makes no distinctions between programs and data, it provides ways to transform code into text and the other way round (namely, the functions `parse` and `deparse`).

Non-standard evaluation and metaprogramming are possible by leveraging delayed evaluation and using functions like `substitute` (that returns the parse tree for an unevaluated expression and at the same time possibly modifies it), `quote` (that simply returns its argument unevaluated) and `eval` (that evaluates an expression in a specified environment). R also provides means for creating embedded domain specific languages (e.g., the formula specification or the “grammar of graphics” of `ggplot2`).¹⁵

To conclude, R being the mixture of different paradigms that it is makes it quite difficult to reason about the code and optimize it. Adding together functional style, object systems, laziness, introspection, dynamic evaluation, computation on the language itself, explicit environment manipulation and more creates a very complex result.

Additionally, R has its semantics defined by its one major implementation (although attempts have been made to formalize the language, e.g., [4]). This further complicates things as there is no formal description of the language.

¹⁵[\[\[about ggplot2\]\]](#)

About RIR

[[Introduce RIR]] RIR is an alternative compiler for the R language.¹⁶ It comes with its own internal representation, an interpreter for its bytecode and an abstract interpretation framework which provides a way to easily implement static analyses on top of the RIR bytecode.

[[history: research project, northeastern? first appearance?]]

RIR acts as a drop-in replacement for the GNU R bytecode compiler. It requires a patched version of GNU R that makes some slight adjustments that allow the standard GNU R expression evaluator function to interface with the RIR bytecode compiler and interpreter. RIR is written in C and C++ and is compiled as a shared library that can be dynamically loaded by R.

The architecture is very similar to GNU R. The compiler is

[[write about rir bytecode]]

[[how is rir bc different]]

[[optimizations, ai framework...]]

[[guards]]

¹⁶Homepage: <https://github.com/reactorlabs/rir>

2.1 Why is RIR slow

[[..., calls to ast interpreter, profiling, optimizations, allocations]]

CHAPTER 3

Improvements

In this chapter I will discuss in detail the changes made to RIR in an attempt to bring it up to speed with GNU R byte-compiled code. The improvements can be subdivided into several categories:

The first group consists of the extensions of the bytecode instruction set itself. Here, instructions are added for some functions that GNU R inlines but were missing in RIR (and thus were compiled as standard calls).

The next section talks about changes made to the RIR compiler. These consist mainly of loop context handling.

Finally, the RIR interpreter loop itself was refactored to run more efficiently, and this is described in the last section of this chapter.

3.1 Instruction set extensions

The GNU R bytecode compiler assumes certain invariants about the code at compile time, as was described in section 1.3.1. Of course, the instruction set of the default compiler reflects this. In fact, having specialized bytecode instructions for specific tasks is where the compiler gets most of its speedups. Specifically, not inlining the primitive R functions of type *special* causes the call mechanism to fall back to the same C routines that the AST interpreter uses, where the expression tree is examined and parts of it are evaluated as needed.

3. IMPROVEMENTS

The first step was to work through the document [6] by the author of the GNU R compiler. Here, all the inlining done by the compiler by default was determined and experiments were carried out to compare the list to RIR, usually by using the disassemblers of both GNU R and RIR and examining the results of compilation of different calls.

The GNU R compiler has four different levels of the *optimize* option (from zero to three) and the amount of inlining performed is directly influenced by this setting. The default value is two, at which the compiler inlines functions in the base packages (including those that are syntactically special or considered core language functions EXAMPLES) that are not shadowed at compile time (by function arguments and local bindings).

The inlining happens in both GUN R and RIR when the compiler sees a function call. Both first try to inline calls to special and builtin functions. If the inlining fails, the compilers fall back to the standard call mechanism. Thus the way to speed things up is to add more special cases that handle code that originally went to the default call.

When adding a new bytecode instruction, several steps have to be performed. First, the instruction has to be added to the instruction list in the *insns.h* header file. The new instruction needs to have a name and also have some properties specified. These are *imm* (the number of immediate arguments that the instruction expects – immediates are inserted directly into the code stream), *pop* and *push* (the number of elements that the instruction removes from and adds to the stack, respectively) and *pure* (a flag that says if the instruction has side-effects, e.g., it can force a promise).

The file with the instructions is intended to be included everywhere that a static list of all instructions (or their properties) is needed. For this reason, all instructions are wrapped in a C preprocessor macro `DEF_INSTR` and this macro has to be defined before including the file. Using this mechanism, one can for instance get an enumeration of all opcodes by defining it to just take the name of the instructions followed by a comma, or a switch over all instructions by putting the names in switch cases.

Second, the instruction has to be manually added at some places in the class that implements the bytecode instruction type and is used throughout the compiler and the analysis framework. This step is mostly mechanical.

After that, the compiler must be taught to use the instruction. This is done during the inlining step. A special case for the function call that the instruction implements has to be added. In this special case, the instruction opcode, together with any other instructions needed (such as the guard instructions explained in chapter 2) are inserted into a code stream. It is also here while the instructions are being added into the stream that constant pool is filled (since the indices of constant pool objects are needed as immediates).

Finally, the instruction itself has to be implemented and added to the dispatching mechanism in the interpreter. This is where the code that executes the instruction is located.

In the following sections, various kinds of added instructions are described.

3.1.1 Relational operators

Originally, RIR only had `<` out of the six usual relational operators. The rest, `>`, `<=`, `>=`, `==` and `!=` were being compiled as calls to the builtin C routines, but were added.

All these instructions behave in the same way as arithmetic binary operators do. They take no immediate arguments, and instead expect their operands to be left for them at the top of the stack. They pop two values off the stack, compute the respective operation, and push one result back. They are not pure.

The speedup of adding these operations as standalone instructions lies in the fact that they are quite often called with “scalar” arguments, and as was mentioned before, R does not have any scalar values, since they are simply boxed in vectors of length one.

If the operands are of a suitable type, and both have a single element, then a fast path can be taken and a call to the standard C builtin avoided. This is possible because internally R uses the C types `int` and `double`, so the actual comparison can be done in plain C.

If not, the instruction falls back to the standard routine that handles vectorized operations, vector recycling, type promotion and also any possible problems (e.g., concerning incompatible operand types). For every instruction, a static variable is also used in this case to cache the builtin function (or rather a pointer to it), so that the lookup (which is an expensive operation) is only performed once.

For the `<` operator, combinations of integer and double scalars had fast paths implemented. This was amended and for all relational operators there is now also a fast path for comparing two logical values.

Furthermore, the relational operators do not need to allocate a new vector for their result, since the R logical objects are singletons.

3.1.2 Unary operators

Unary operators are in principle the same as binary. If their operand has length one and appropriate type, a fast path can be added.

R has two arithmetic unary operators, `+` and `-`, and logical negation `!`. These are all rather straightforward, they do not have immediates, pop one value and push one value.

The logical negation has a fast path for logical scalars in addition to numeric.

3.1.3 The colon operator

The colon operator `:` in R provides a convenient way to generate sequences. It is used very often, notably in `for` loops as an integer control sequence for the loop variable. The values it generates can be both increasing and decreasing, and they differ by one. If the starting value is an integer, then the vector is also integer.

The colon instruction was added that, similarly to arithmetic operators, takes no immediates, expects two operands on the top of the stack and pushed back the resulting sequence object. It is impure. Adding it to the compiler was actually the same as adding an arithmetic binary operator. The instruction itself adds a fast path for combinations of integer and double operands (the doubles

apply only if they represent an integer up to a rounding error). This fast path allocates an integer vector of appropriate length and fills it with the sequence values.

3.1.4 Superassignment

Superassignment operator `<<-` differs from normal assignment in that it works in an enclosing environment. Thus, the local environment where superassignment occurs is skipped during the binding lookup. For simple assignment of the form `x <- y`, that is all there really is.

The semantics of a complex subset assignment are a bit more complicated, as is shown in listing 3.1 (taken from [5]), because it is not a matter of simply creating or changing a binding, but a part of a binding's object has to be extracted and modified first. This model applies recursively for still more complex assignments.

The target object is looked up and stored into a variable `*tmp*` (it is not recommended to use this name for anything in user code, since, as a side-effect, this will overwrite and then delete it). Then, a function name is constructed from the left-hand side call expression by appending an assignment arrow. The resulting name must refer to a function that takes the same arguments as the original one as well as an additional value argument. This function is called, the right-hand side expression is passed as the value, and its result is stored to the target. Then the temporary binding is removed.

For superassignment the same principles apply, however, the target binding (and only the target binding) is looked up in an enclosing environment of the expression.

```
# The result of this command...
x[3:5] <- 13:15
# ... is as if the following had been executed
*tmp* <- x
x <- "[<-"(*tmp*, 3:5, value=13:15)
rm(*tmp*)
```

Listing 3.1: Complex subset assignment

3. IMPROVEMENTS

Two new bytecode instructions were added for handling the superassignment semantics of looking up bindings. The first is for loading a symbol. It takes one immediate argument, an index into the constant pool where it finds the symbol to look up. It does not pop anything from the stack but pushes one object, the value of the binding. It is not pure. The second is symmetrical to the first, it takes an immediate constant pool index, pops one object, does not push anything and is pure.

These new bytecodes were added to the compiler at a point where it inlines normal assignment and subset assignment.

In the instructions themselves, the environment for looking up bindings gets replaced with its enclosing environment.

3.2 Compiler modifications

[[loop contexts, loop refactoring, bc cleanup]]

```
function(n) {  
  repeat {  
    if (n <= 0) break  
    n <- n - 1  
  }  
}
```

Listing 3.2: Safe **break**

```
function(n) {  
  repeat {  
    foo(if (n <= 0) break else 3)  
    n <- n - 1  
  }  
}
```


Listing 3.3: Context for **break** required

3.3 Interpreter refactoring

As it turned out, the biggest speedup was gained by refactoring the RIR bytecode interpreter. Originally, the interpreter was quite straightforward. The main evaluator function `evalRirCode` contained in its core an infinite loop, and in its body there was a huge switch statement with one case for each bytecode instruction.


In the cases there was a call to a function that implemented the instruction. These function were defined with the C `inline` and `static` modifiers, and to make sure the compiler really inlined them, the GNU extension attribute `__attribute__((always_inline))` was also specified.

However, something



Listing 3.4

```
/todo[register]
```

Evaluation

[[discussion of results, add figures]] [[discuss interesting point in measurements - naive nbody and threading etc.]]

[[plot times without warmup]]

[[plot how times change with successive invocations]]

4. EVALUATION

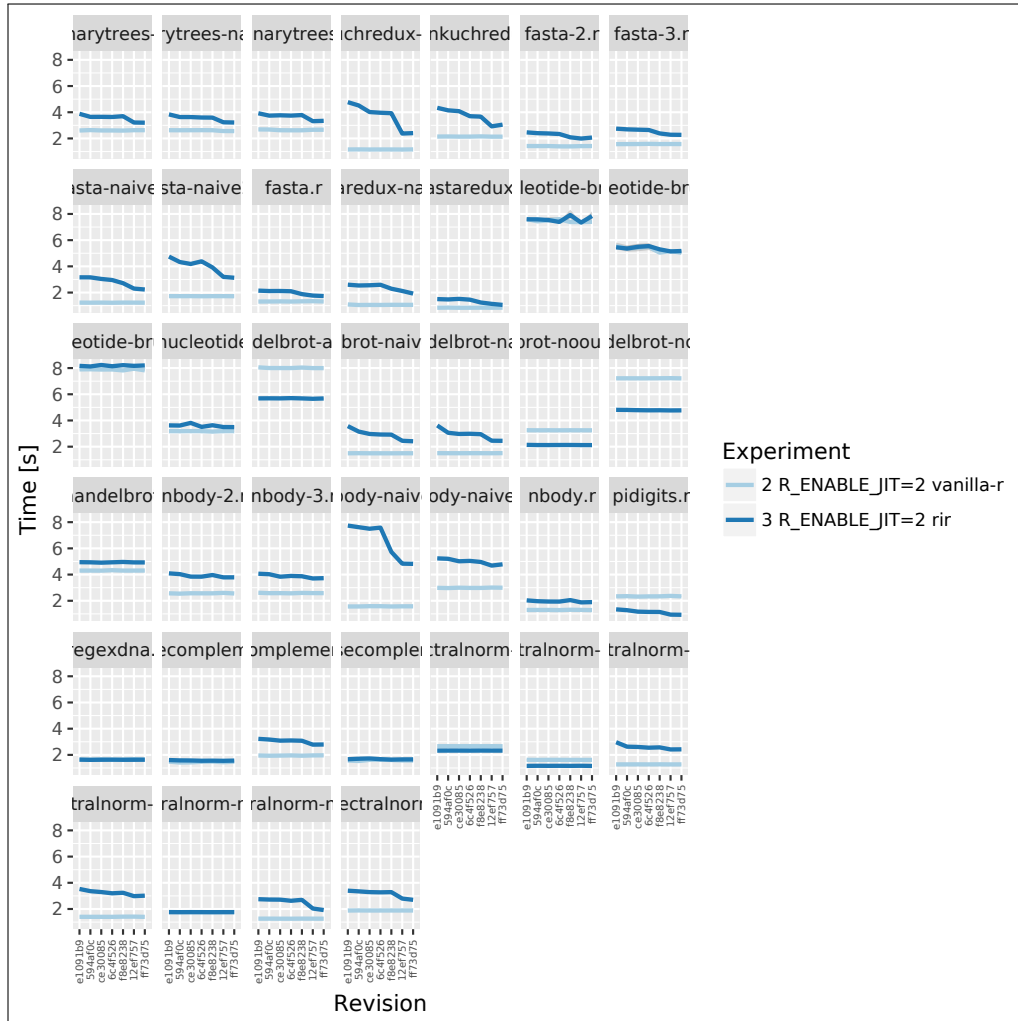


Figure 4.1: History of running times




Figure 4.2: Overview of the speedup vs. GNU R

4. EVALUATION




Figure 4.3: History of average speedup vs. GNU R

Conclusion

[[conclusion, future work, related work, fails - promises of consts, stoke etc.]]

Bibliography

1. BURNS, Patrick. *The R Inferno* [online]. 2011 [visited on Apr. 30, 2017]. Available from: <http://www.burns-stat.com/documents/books/the-r-inferno/>.
2. DALGAARD, Peter. *The R-announce Archives: R 2.13.0 is released* [online]. 2011 [visited on Apr. 30, 2017]. Available from: <https://stat.ethz.ch/pipermail/r-announce/2011/000538.html>.
3. DALGAARD, Peter. *The R-announce Archives: R 3.4.0 is released* [online]. 2017 [visited on Apr. 30, 2017]. Available from: <https://stat.ethz.ch/pipermail/r-announce/2017/000612.html>.
4. MORANDAT, Floréal; HILL, Brandon; OSVALD, Leo; VITEK, Jan. *Evaluating the Design of the R Language: Objects and Functions For Data Analysis* [online]. 2012 [visited on Apr. 30, 2017]. Available from: <http://r.cs.purdue.edu/pub/ecoop12.pdf>.
5. R CORE TEAM. *R Language Definition: Subset assignment* [online]. 2017 [visited on Apr. 30, 2017]. Available from: <https://cran.r-project.org/doc/manuals/r-release/R-lang.html#Subset-assignment>.
6. TIERNEY, Luke. *A Byte Code Compiler for R* [online]. 2016 [visited on Apr. 30, 2017]. Available from: <http://homepage.divms.uiowa.edu/~luke/R/compiler/compiler.pdf>.
7. WICKHAM, Hadley. *Advanced R* [online]. 2014 [visited on Apr. 30, 2017]. Available from: <http://adv-r.had.co.nz/>.

Acronyms

API	Application programming interface
AST	Abstract syntax tree
BC	Bytecode
CLI	Command line interface
CRAN	The Comprehensive R Archive Network
GNU	GNU's Not Unix!
IDE	Integrated development environment
JIT	Just-in-time compilation
REPL	Read-eval-print loop
VM	Virtual machine

Contents of the enclosed CD

[[contents of cd]] After this fourth paragraph, we start a new paragraph sequence. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.